



A DECENTRALISED PEER-PREDICTION MARKET

CS907 DISSERTATION PROJECT - INTERIM REPORT

Thomas Archbold
1602581

Department of Computer Science
University of Warwick

July 30, 2020

1 Introduction

Prediction markets are exchange-traded markets¹ which allow users to trade on the outcomes of future events as opposed to traditional financial instruments. Users participate by placing bets and buying or selling shares in various markets. Since users stake their own money the market prices should indicate the true beliefs of the userbase and the perceived likelihood of events occurring. Different users will have different beliefs and knowledge informing their decisions, hence a prediction market provides a means of aggregating information on topics of interest using “the wisdom of the crowd”. Securities in these markets are usually traded between \$0 and \$1, and for binary events a market will typically pay out \$1 for every share held for a positive outcome, and \$0 otherwise.

Traditional prediction markets are centralised in the sense that the system provides the bets for the users to trade in, and traders simply choose at what price point and quantity of shares to participate. Since the system knows all possible bets beforehand it is then easy for it to be responsible for determining the outcomes of these markets and paying out the winnings accordingly. There are two issues with this centralised approach: firstly, it restricts the types of bets that can be made, since they must be explicitly offered by the market maker; secondly, it operates on trust – there is nothing to stop the central market maker from manipulating the system for their own gain. We aim for a system that avoids both of these issues.

In this project we implement a *decentralised* prediction market in which it is up to the userbase itself to define the markets and determine the outcomes of these events – these outcomes are decided upon by consensus among a group of users, known as arbiters. This removes the need for a trusted centre, however with no central moderator bets may become ambiguous or their outcomes subjective, and arbiters may still attempt to manipulate the outcome of the market for their own gain by submitting false reports. It is also important that users continue to act according to their true beliefs. We base our design on the incentive compatible outcome determination mechanism proposed by Freeman, Lahaie, and Pennock [7].

The rest of the report is structured as follows: in Section 2 we give an overview of the current literature on prediction markets, and in particular decentralised markets. In Section 3 we outline the specific goals of the project and briefly detail the current progress we have made. In Section 4 we discuss the high-level system design, tools we use, and implementation details. In Section 5 we detail the approach taken to the project’s development, ethical considerations, and scheduling issues, and lay out plans for next tasks to complete. Finally, in Section 6 we reflect on successes and failures of the project so far.

2 Background

Prediction markets receive considerable attention within the algorithmic game theory literature. They have been used to predict the outcomes of a range of events, particularly in business, sports, and politics [15, 6]. Nay, Van der Linden, and Gilligan [13] analyse prediction markets empirically to study how market parameters may affect an agent’s beliefs on the real-world outcome, making the argument for a prediction market for betting on the cause of global warming to show how increased agent participation can cause public belief to shift towards the “correct” cause. They are therefore powerful not only in aggregating public opinion, but also in their ability to change it.

In the centralised setting Kroer, Dudík, Lahaie, and Balakrishnan [11] introduce a cost-based market maker in which all bets are bought and sold to the market, rather than the traders, which sidesteps the problem of low liquidity in traditional markets. Another way to avoid low liquidity is to use a market scoring rule: Liu, Wang, and Chen [12] introduce Surrogate Scoring Rules to elicit private probabilistic beliefs with access only to agents’ reports, which is similar to our setting. Their mechanisms do not rely on a ground truth to quantify the quality of the elicited information, and could hence be particularly applicable to decentralised prediction markets. Freeman et al. [7] take such a decentralised approach and introduce a mechanism in which market outcomes are crowdsourced by a group of arbiters. They derive the conditions under which arbiters are incentivised to act truthfully, even if an arbiter reports on a market where they themselves hold a stake, and theirs is the paper upon which we base our design. Several other decentralised markets exist: Peterson et al. [14] study the setting to implement the decentralised oracle behind the prediction market *Augur* [2], which uses a reputation system to protect against manipulation. Similar markets based on cryptocurrencies exist in *Gnosis* [3] and *Hivemind* [4], though it has not been shown that any of these achieve any theoretical guarantees of truthfulness.

¹Markets in which all transactions are routed through a central source.

3 Goals

The overall goal of the project is to implement a decentralised prediction market in which users may act as arbiters in markets in which they themselves are stakeholders. This will be achieved by the following essential objectives:

- create a web application where users may specify bets to trade on
- implement a trading mechanism allowing users to buy and sell shares in user-made securities
- crowdsource outcome determination using reports from arbiters who may hold positions in the markets
- incentivise truthful reporting from the arbiters

These goals will largely be achieved by implementing the mechanism laid out by Freeman et al [7]. The first three goals cover core functionality of any decentralised prediction market, while the fourth is concerned with tuning system parameters, before and during execution, in order to ensure that users do not manipulate the mechanism.

Other features that we look to implement that are not essential to the running of the system but highly desirable for user experience include asynchronous communication with the server to display up-to-date pricing information without a page refresh, and the automated closing of markets. Indeed, while their current omission does not necessarily hamper the functionality of the market itself, they should be implemented in order for the system to run smoothly and independently.

The following additional features could be implemented, although they are not the focus of the project or are unlikely to be achieved in the remaining time frame. They include removing the need to explicitly ask for signal beliefs from arbiters; the option to sort markets by categories; displaying user trade histories; and sanitising database inputs.

3.1 Progress

The project has so far achieved the first three goals, meaning the core functionalities of a prediction market are implemented. Users are able to create markets by specifying a wager and a deadline by which the outcome will be known, and may then trade in these markets. Buying shares pushes the price up while selling them pushes it down, both of which aggregate the perceived likelihood of the event. Profits can be made in the markets in the traditional way, by buying low and selling high (before the deadline has passed), or by maintaining the position past the deadline and receiving a payout based on the outcome of the wager.

We have also implemented crowdsourced outcome determination, in which we resolve the outcomes of bets by collecting reports from users. These users offer their opinion on the market's outcome, a "yes" or a "no", as well as some extra information regarding their belief of the accuracy of their signal, and are paid a certain amount of money as a reward for helping to resolve the market. The payout per share is then set to the fraction of arbiters reporting a positive result, and all users with a position are paid out, or made to buy back shares, accordingly.

Currently, although we compute many of the values with the correct method, it can not yet be said that we achieve truth-telling in the arbitration stage since many of the parameters on which these values depend have been set arbitrarily and are not yet tuned properly.

4 Design

4.1 Tools

The project is implemented in Common Lisp and the code has been developed and tested within the Steel Bank Common Lisp (SBCL) compiler and runtime system. Code version control has been achieved with Git and Github. Writing the web application has required the use of several libraries available from the library manager Quicklisp, specifically:

- Hunchentoot
- CL-WHO
- Mito

- SXQL
- Parenscrip
- Smackjack

Hunchentoot provides the environment on which we host the server. Most importantly it provides automatic session handling, allowing for multiple users to be logged in at once, and easy access of GET and POST parameters, enabling interaction via HTML forms. To generate the webpages, we use CL-WHO, which converts Lisp statements into strings of valid HTML. Defining webpages while remaining in the Lisp environment means we may use Lisp’s macro system to build abstractions for both defining structure and processing data in one interface.

Mito and SXQL provide the ability to connect to and interact with a Relational Database Management System (RDBMS). We use MySQL, though this choice is largely immaterial given our simple requirements of the database.

Parenscrip incorporates Javascript into the site with the goal of improving user experience. Currently it is only used to ensure that all necessary fields during market creation, trading, and market resolution are non-empty to avoid sending incomplete data to the server. It will be used to a greater degree in the future to ensure responsiveness: all information displayed to the user must be current to ensure that users are interacting with an up-to-date state of the system. We therefore plan to make more use of Parenscrip and Smackjack, an AJAX library for Lisp, to allow for asynchronous interaction with the server. This will include, for example, continuously updating a stock’s price or calculating the cost of a transaction without a page refresh, and stronger client-side validation.

4.2 System Overview

We implement the mechanism for peer prediction introduced by Freeman et al [7]. In this section we will introduce the main ideas they present and give an overview of the mechanism itself.

We are interested in setting up a prediction market for the outcome of the binary event (random variable) $X \in \{0, 1\}$. The terms “market” and “security” are used interchangeably to refer to the entity comprising a wager (e.g. “Arsenal will beat Tottenham”) and a deadline – these two pieces of information are all we need to represent event X . The mechanism is divided into two main stages: the market stage, where users may buy and sell shares in the securities whose deadlines have not yet passed; and the arbitration stage, where a subset of the users report on the outcome of the security and we compute the payout price per share held. A key feature of this mechanism is that users may act as arbiters in markets in which they themselves hold shares.

Since we rely heavily on user participation for the mechanism to run correctly, it is important that users act in the desired manner. This mechanism incentivises truthful reporting in the arbitration stage, whereby it is in a user’s best interests to report what they truly believe a market’s outcome to be and not attempt to manipulate the system, even if the user is a stakeholder in the market. It also allows us to achieve certain guarantees, such as a bound on the amount we must pay to arbiters to reward them for participation, which serves as a guide when setting system parameters.

4.2.1 Market Stage

The market stage allows users to create markets for any event they wish and trade shares in these markets. Since we crowdsource outcome determination, there is no restriction placed on the types of bets users may make other than that their result is a “yes” or a “no”. Ambiguous bets are allowed, though likely to suffer in the arbitration stage, since users may have different interpretations on the outcome.

In order to create a market for a user-specified event, we use a *market scoring rule (MSR)*, which is a means of assigning a probability to a set of mutually outcomes. After the wager and deadline have been specified, the MSR simply takes into account the number of shares bought and sold and returns a probability $p \in [0, 1]$. This describes the perceived likelihood amongst users of the event having a positive outcome. We can also use this to set the share price of the security. Using a MSR is different from a traditional market in that there is not a fixed number of shares in circulation: instead, buying into a market increases the total number of shares and selling decreases it. Let q denote the total number of shares of a given security. An agent wishing to buy $q' - q$ shares (i.e. increasing the total number of shares to q') will pay $C(q') - C(q)$, for our choice of convex, differentiable, monotonically increasing scoring rule C . We may tune the behaviour of this cost function by using a liquidity parameter b , such

that $C_b(q) := b \cdot C(q/b)$, which controls the responsiveness of C . A lower value of b means that the share price changes more quickly for fewer shares bought, and vice versa.

The market stage implements trading fees in order to raise the funds necessary to pay stakeholders when the event's outcome is realised, and to reward users for participation in the arbitration stage. Buying shares pushes the share price p towards \$1, while selling shares pushes it towards \$0, and the fee can be interpreted as a fee on the worst-case loss that an agent incurs. For a fixed parameter f , a “buy” transaction that pushes the share price to p incurs an additional charge of fp , while a “sell” transaction incurs a charge of $f(1 - p)$. Transactions are only charged a fee when the user increases their risk: if they are simply liquidating their position (selling shares they own or buying back shares they have sold) then no fee is charged. Users may trade in the market as long as they have enough funds to make the transaction and the deadline for the event has not passed. After the deadline has expired the users' positions are final and we move on to determining the outcome of the event in the arbitration stage.

4.2.2 Arbitration Stage

The arbitration stage is concerned with determining the perceived outcome of event X by taking reports from a subset of the users in the system, known as the arbiters. Each arbiter i receives a (private) signal $x_i \in \{0, 1\}$ that tells them the result of the event – this is analogous to reading the news, watching the match, or even hearing about it from a friend. They then submit a report $\hat{x}_i \in \{0, 1\}$ to the system that tells us what they believe the outcome to be. Note that since the signal they receive is private, we have no way to know whether the user is reporting what they truly believe, or whether they are trying to manipulate the system for their own gain. Arbiters are then paired randomly, and paired arbiters i and j are paid a reward of $u(\hat{x}_i, \hat{x}_j)$ according to the “1/prior with midpoint” mechanism, which we will detail below.

We can now determine the outcome of the market. In a traditional prediction market, users are paid \$1 for each share owned (or they must buy shares back at \$1 if they have shorted the security) if the event has a positive outcome, and \$0 otherwise. Hence if an event is almost certain to occur, more users will buy shares than sell them, as they expect they will be paid out for holding shares in the market when the outcome is realised. This will push the share price towards \$1. Similarly, if an event is believed to be unlikely, more users will sell shares than buy them and the price will approach \$0. This prevents arbitrarily large profits being made for risk-free bets. In our market the outcome of an event is the random variable $\hat{X} \in [0, 1]$ which is set to the proportion of arbiters that submit a report of $\hat{x}_i = 1$. Stakeholders are then paid out in the usual way.

4.2.3 1/prior mechanism

We use the 1/prior payment mechanism to reward users for submitting reports on the outcomes of events, with a modification to incentivise truthful reporting. The 1/prior payment mechanism was conceived by Jurca and Faltings [9, 10] as a means of rewarding arbiters for participation in opinion polls, and Witkowski [16] generalises this mechanism to pay out different amounts depending on the signals reported by paired arbiters. For paired arbiters i and j with reports \hat{x}_i and \hat{x}_j , the 1/prior mechanism pays i and j as a reward for reporting on the market outcome:

$$u(\hat{x}_i, \hat{x}_j) = \begin{cases} k\mu & \text{if } \hat{x}_i = \hat{x}_j = 0 \\ k(1 - \mu) & \text{if } \hat{x}_i = \hat{x}_j = 1 \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

in which k is a parameter and μ is the (common) prior belief that $X = 1$. A suitable value to use for μ would be the closing price for the market: if users feel an event is likely to happen the share price will be pushed towards \$1, while if they feel it is unlikely it will be pushed towards \$0.

The modification to the 1/prior mechanism is simple though requires two extra values to be computed. Let μ_1^i be the probability that, given that agent i receives a positive signal, another randomly chosen agent also receives a positive signal, and let μ_0^i be the probability that, given that agent i receives a negative signal, another randomly chosen agent receives a positive signal. We require the “update” probabilities μ_1, μ_0 to be common across all agents, which we can achieve by simply taking:

$$\begin{aligned} \mu_1 &:= \min_i \mu_1^i \\ \mu_0 &:= \max_i \mu_0^i \end{aligned} \quad (2)$$

The modified payment method we use, the “1/prior with midpoint” mechanism, is now simply Equation 1 with μ replaced by $(\mu_1 + \mu_0)/2$. This guarantees the incentives for arbiters are the same regardless of their signal.

4.3 Implementation

As discussed, we have currently implemented both stages of the mechanism but are yet to tune the parameters required to achieve the truth-telling behaviour we desire. The system comprises four independent parts, each implementing a separate area of functionality, as follows:

1. Database
2. Trading
3. Arbitration
4. Server

4.3.1 Database

The database is implemented using the Mito library provided by Quicklisp. This is an object relational mapper that provides support for MySQL, PostgreSQL, and SQLite3. We can therefore define and interact with tables while staying completely within the Lisp ecosystem, leading to a quicker and more flexible development cycle.

We use three tables: **USER**, which stores all users in the system and their remaining budget; **SECURITY**, which stores the wager, deadline, number of shares, and outcome of every user-created security; and **USER-SECURITY**, which is responsible for the many-to-many relationship between users and securities (i.e., each user may own shares in many securities). The latter allows us to store a user’s position within a market as well as how they report on it during the arbitration stage, if they indeed decide to do so. The tables contain the following columns:

User	Name	Budget
-------------	------	--------

Security	Bet	Shares	Deadline	Outcome
-----------------	-----	--------	----------	---------

User-security	User	Security	Shares	Report	Positive Belief	Negative Belief
----------------------	------	----------	--------	--------	-----------------	-----------------

An entry in the **USER** table consists of a username and a budget. There is currently no requirement to enter a password to log in to the system – this has been done to speed up the testing process, and will be an easy addition in the future.

The columns in the **SECURITY** table contain all information to define a market in which users may trade: “bet” holds the wager that is being proposed, “shares” holds the number of shares which have been bought and sold, “deadline” holds the earliest date-time by which users can receive a signal on the wager’s outcome, and “outcome” holds the payout price per share of the market, to pay out stakeholders.

An entry in the **USER-SECURITY** table stores the relationship between users and securities, whether this is the number of shares a user holds in a security, an arbiter’s report on a security, or both. The “user” column holds the user ID, “security” holds the security ID, and the number of shares owned and their report (or NULL) are stored in the corresponding columns. The columns “positive belief” and “negative belief” store a user’s estimations that, given they the event has a positive (respectively negative) outcome, they receive a positive signal. The manner in which this is used is detailed in Section 4.3.3.

Defining such tables using Mito is as simple as using the `deftable` macro, and is syntactically similar to defining a regular struct within Common Lisp. Listing 1 shows how we define the columns and their associated datatypes. The macro defines default accessors², the slots `created_at`, `updated_at`, and a primary key `id` if none is specified. Insertion is similarly straightforward: to insert a new entry into a table we create an instance of the class that is implicitly defined when calling `deftable`, then call `create-dao`. To retrieve records from the database we use either `select-dao`, which returns all records matching the conditions, or `find-dao`, which returns only the first.

²Functions for accessing members of a struct.

```

1 (deftable security ()
2   ((bet :col-type :text)
3     (shares :initform 0
4       :col-type :integer)
5     (deadline :col-type :datetime)
6     (outcome :col-type (or :double :null))))

```

Listing 1: Defining a table in Mito

4.3.2 Trading

Trading allows users to create arbitrary bets on binary events and then trade shares in them. It also implements the trading fees on risk transactions required to raise funds to reward arbiters and pay out stakeholders. To calculate the trading fee, we first need to know whether the user is increasing their risk (a “risk transaction”) or they are simply liquidating a position they already hold, and secondly whether they are buying or selling shares, since these incur different charges.

To create markets, quote their share prices, and compute the cost of a transaction we use a market scoring rule. We use the commonly used Logarithmic Market Scoring Rule (LMSR), created by Robin Hanson [8]. As we have mentioned, in any MSR an agent wishing to purchase $q' - q$ shares of a given security must pay $C(q') - C(q)$ for convex, differentiable, monotonically increasing function C . In the case of LMSR, for current number of shares q and liquidity parameter b , this function is:

$$C_b(q) = b \log(1 + e^{q/b}) \quad (3)$$

This function also accommodates selling shares, in which case $q' < q$. Since a user must specify a desired amount of shares to buy or sell before we can determine how much to charge them, we cannot use this function to directly quote a share price. Instead, we can calculate the cost of the transaction were the user to purchase a miniscule quantity of shares and quote this as the share price. This is simply the derivative of Equation 3 and is given by:

$$p_b(q) = \frac{e^{q/b}}{1 + e^{q/b}} \quad (4)$$

Although Equation 4 quotes the share price, it is not accurate to use it to calculate how much an agent should pay, since the agent’s participation in the market will immediately change this price. It does, however, provide a useful means of describing to the user the perceived likelihood of the event. For example for some security, at $q = 0$ an equal number of shares have been bought as have been sold, meaning the userbase as a whole is equally unsure about whether the market will have a positive or negative outcome. The quoted share price will be $p(0) = \$0.50$. Similarly, suppose $b = 10$ and $q = 20$, meaning twenty more shares have been bought than sold in the market. This would yield a share price of $p_{10}(20) \approx \$0.88$. Users feel more confident that the event will have a positive outcome, since more people have bought shares, giving a higher share price. As mentioned in Section 4.2.3, we interpret the closing price of the market as the probability of a positive outcome, as it approximates how sure users feel about the event’s outcome.

home about portfolio logout

User: tom Funds: 75.2467

Create Position

Market	The weather will be hot on Friday	
Deadline	2020-07-31 00:00	
Share price	0.5000	
Initial Position	<input type="text" value="4"/>	shares
Projected Cost	TODO (update asynchronously)	

Buy shares

Figure 1: Creating a market for the wager, “The weather will be hot on Friday”

4.3.3 Arbitration

The arbitration stage of the mechanism seeks to resolve markets whose deadlines have passed by gathering reports from arbiters and paying out, or demanding money from, stakeholders. Since we will eventually be able to accommodate for users to act as arbiters in market in which they themselves hold a position, we allow any user to opt in to become an arbiter in an unresolved market. Figure 2 shows the dashboard as it appears to a logged in user, in which they are presented with all the currently active and unresolved markets.

The dashboard for a user named 'tom' with funds of 73.0180. It features a navigation bar with links to home, about, portfolio, and logout. The main content is divided into three sections:

- Active Markets:** A table with columns 'Bet', 'Deadline', and 'Price (\$)'. It lists two active markets:

Bet	Deadline	Price (\$)
"Brentford will beat Swansea City"	21:30 29-07-2020	0.7311
"The weather will be hot on Friday"	00:00 31-07-2020	0.5987

 Each row has a 'Trade' button.
- Unresolved Markets:** A table with columns 'Bet', 'Expired on', and 'Closing price'. It lists two unresolved markets:

Bet	Expired on	Closing price
"Watford will get relegated"	16:30 26-07-2020	0.8808
"Elon Musk will make a silly tweet"	13:00 28-07-2020	0.6682

 Each row has 'Report Outcome' and 'Close Market' buttons.
- Create a Market:** A form with fields for 'Bet', 'Deadline' (with a date and time picker), and a 'Create market' button.

Figure 2: A user's dashboard, where they can see all active and unresolved markets

Once the user has opted in to arbitration they are taken to a form to submit their report on the outcome, as well as estimations on their "Positive Belief" and "Negative Belief". These are estimations of probabilities that, given that the event *actually* had a positive (respectively negative) outcome, the user received a positive signal. This is important as we need to take into account signal noise, given that wagers can be subjective: while the positive and negative signal beliefs will be nearer to 1 and 0, respectively, for an event such as the winner of a football match, since it will likely be covered by many media outlets all reporting the true result, this may not necessarily be the case where the outcome is more a matter of opinion. The weakness of this approach is discussed in Section 6.

The 'Resolve Security' form for the market "Elon Musk will make a silly tweet" with a deadline of 13:00 28-07-2020. It includes a 'Market Outcome' section with radio buttons for 'Yes' (selected) and 'No'. Below this are two input fields for 'Positive Belief' (set to 80%) and 'Negative Belief' (set to 40%). A 'Submit' button is at the bottom.

Figure 3: The interface for submitting a report as an arbiter

These signal beliefs are used to compute the update probabilities μ_1 and μ_0 which are used in the 1/prior-with-midpoint mechanism. Given the closing price μ and the positive and negative signal beliefs $Pr[x_i = 1|X = 1]$ and $Pr[x_i = 1|X = 0]$ for each arbiter i , we can compute the positive update μ_1^i for each i and randomly chosen j as follows:

$$\begin{aligned} \mu_1^i &= Pr[x_j = 1|x_i = 1] \\ &= Pr[x_j = 1|X = 0] \cdot Pr[X = 0|x_i = 1] + Pr[x_j = 1|X = 1] \cdot Pr[X = 1|x_i = 1] \end{aligned} \quad (5)$$

We use the same approach to compute the negative update μ_0^i for each i , and hence we may now also compute the common updates μ_1 and μ_0 to pay arbiters the correct reward for arbitration. Once done,

we set the outcome of the market to the fraction of arbiters that reported a positive outcome, paying users who hold shares and taking money from users that sold shares short.

4.3.4 Server

The code contained in `server.lisp` is responsible both for setting up and maintaining the Hunchentoot web server and defining the various pages and forms with which users interact. Hunchentoot makes interaction with the web server straightforward. We initialise the server by creating an `easy-acceptor`, and once we have defined our webpages using CL-WHO we simply push them to the dispatch table using `create-prefix-dispatcher` so that they may be accessed. Hunchentoot also provides us with automatic session handling, meaning we do not have to worry about the details of logging in users on different machines at the same time: we simply define the symbol `session-user` in the appropriate session-handling data structure, then set this to the logged in user. We can then access this using (`session-value 'session-user`) to display different material on the page and interact with the database according to the current user. Finally, we use Hunchentoot to access the GET and POST parameters users send to the server via forms with a call to the library's `parameter` function.

Although there is little special to CL-WHO compared to other HTML-generating libraries, it is useful to remain in the Lisp environment to define webpages since we may use its powerful macro system. We use it to define a template for a standard page, meaning we only need to specify the elements that make the page unique and giving the web application a consistent style. This also enables us to define pages and add the corresponding HTML generating function to the dispatch table all within a single interface, hiding the unnecessary details. Listing 2 shows how we are able to define both the webpage, generate its content, and push it to the dispatch table in one.

The code in `server.lisp` draws together the interfaces from other areas of the system and actually executes the mechanism. This involves taking user input to create custom markets, quote share prices, and collect and distribute payments for the various interactions users may have with the market.

```
1 (defmacro define-url-fn ((name) &body body)
2   '(progn
3     (defun ,name () ,@body)
4     (push (create-prefix-dispatcher ,(format NIL "~(~a~)" name) ',name)
5         *dispatch-table*))
```

Listing 2: Defining webpages

5 Project Management

5.1 Methodology

The project has been developed incrementally, with a focus on integrating new functionality completely before progressing to new features. This approach is well-suited to this project's design: since it comprises of mainly four separate areas which are drawn together at the end, it is possible to focus on implementing a feature within one area without it affecting the rest. As a result, testing has been performed throughout and ensures that a newer version of the project is never worse than its predecessor. Using Git and Github has been helpful in this regard, providing cloud storage and the ability to roll back to previous versions if the current one is broken by a new feature.

5.2 Ethics

There is little ethical consideration required for the development of this project. All development and testing has been done independently, and all resources used to implement the system is available freely. Testing has been performed externally only to small extent, and even then only informally through asking of colleagues' opinions. There is one ethical issue that faces prediction markets in general, however, and the one we implement is no exception. When prediction markets operate, directly or indirectly, with real money, they run the risk of devolving into "assassination markets" [1, 5]. This refers to the incentive that people may have to act in a way that changes the "natural" outcome of an event. In the extreme case, a bet may be made on the death of a high-profile individual and an assassin could stand to make a large by participating in the market and ensuring that the individual dies on a given date. A way to sidestep this issue is to use virtual money with no real-world value, and this is the approach we take.

5.3 Scheduling

There have been few issues regarding the scheduling of the project, though it has benefited from a slight rearrangement of the initial timetable. Figure 4 shows the schedule as it was planned at the time of the presentation, towards the start of the project’s initial development. Tasks have also taken slightly longer to implement, and a revised version is given in Figure 5, which details both how time on the project was spent and how we now plan to use the remaining time.




Figure 4: Project timetable as it was initially planned

Firstly, although an early demonstration of the market and arbitration stages was complete for the presentation, this did not yet use any persistent storage or actual users. Therefore, to get these stages working as they would in the final project, the website and database interface first needed implementing. This has taken the majority of the time since the presentation, and includes defining the various pages and forms with which the users interact with the markets, as well as how we interact with the database to register, log in, and pay users. As we will cover in the next section, there are only three main areas of functionality left to target: asynchronous communication with the server via AJAX calls, automation of the system regarding closing markets once enough arbiters have submitted reports, and most importantly parameter tuning. The latter has been given a large block of time to complete since there are many values on which certain parameters depend, and it is imagined that calculating these accurately will require care.

Secondly, the time dedicated to writing the interim report has been extended by two weeks to reflect the extension to its deadline. Although it was not initially required, the extra time that could be diverted towards developing the systems while the writing of the report was delayed was useful in providing more material to discuss. This was also caused in part by more focus than anticipated being paid towards exam preparation.

5.4 Next Steps

The next task is to tune the parameters of the system to ensure truthful reporting, since this is the entire reason for implementing the mechanism. This includes setting transaction fee f to the appropriate percentage and the payment parameter k . It will also be important to fully automate the system, meaning the server automatically decides when to close markets: this is currently done by prompting the system to compute market outcomes, which is insufficient. The final important feature to add is asynchronous server communication – although it does not necessarily add new functionality to the core of the prediction market, it would dramatically increase usability.

We plan to dedicate a significant portion of August to the writing of the dissertation, to spread the workload more evenly and to allow features that are developed later more time to be added into the report. We hope to get feedback on a draft of the dissertation three weeks before the final deadline so that the following weeks can be spent refining the material, and for more versions to be produced in the remaining time. This means we intend to get all required goals working well before the end of the month, to ensure we are not making considerable changes shortly before the deadline.

	04/20	04/27	05/04	05/11	05/18	05/25	06/01	06/08	06/15	06/22	06/29	07/06	07/13	07/20	07/27	08/03	08/10	08/17	08/24	08/31	09/07
Exam period																					
Presentation																					
Interim Report																					
Final Report																					
Market Stage																					
Arbitration Stage																					
Database design																					
Web interface																					
Asynchronous updates																					
Automation																					
Parameter Tuning																					
Prediction Markets literature																					

Figure 5: Revised project timetable

6 Evaluation

6.1 Successes

One main and obvious success of the project so far is that it successfully implements a prediction market. This has required learning about the associated game theory literature on mechanism design and scoring rules markets, and although the project has taken a different shape since its proposal, a better working knowledge on the current landscape of algorithmic game theory was always a goal and has been achieved. An important point raised in the presentation was that the project is not necessarily interesting because it allows users to trade on predictions, but more so because it allows users to create their own markets while remaining robust to manipulation, hence the game theoretic, rather than strictly economic, approach.

The project has required learning the Lisp programming language having had no prior experience in it. Writing custom macros to abstract away unnecessary details in database interaction and HTML generation has provided particular satisfaction.

6.2 Improvements

While there have been no particular failures of the project so far, there are areas of it which can be improved. There is an issue of efficiency in database access caused by an unfamiliarity with the Mito library. It is likely, however, that we will not be able to rectify this in time, as it currently causes no performance issues and there are more important tasks taking priority.

The frequency of supervisor meetings should be made more consistent as we approach the final stages of the project. Until this point this has not been an issue and was anticipated over the summer examination period, however as we begin work on the final report a more steady opportunity to discuss problems and receive feedback will become more useful. Although the irregularity of meetings can be attributed in part to the remote working situation, the facilities exist for meetings to be held remotely and should be used to a greater degree as we approach the project's end.

There are also design aspects of the system that are worth improving. The requirement for users to supply their estimates of signal accuracy is an issue: it is unreasonable to ask for such estimations, and also opens up the opportunity for users to manipulate the mechanism, which is the very behaviour we wished to avoid. This is a weakness of the mechanism from Freeman et al. [7], and the issue arises as the signal posteriors are assumed to be known by the system. In practice this is not possible since we have no way of knowing the nature of where users get their signals. This problem could be alleviated by calculating these probabilities based on past reporting and market outcome histories, removing the interaction required from the user. It would also be worth exploring the use of the Surrogate Scoring Rules introduced by Liu et al. [12], given that they do not need access to a ground truth to score predictions. This could be particularly useful since a minor issue of the mechanism we implement is that it is somewhat aged. This is not of much concern, since no similar prediction market exists in a practical implementation, however it would be interesting to improve on its shortcomings using results and approaches from more recent works.

References

- [1] *Assassination market*. https://en.wikipedia.org/wiki/Assassination_market. Accessed:

2020-06-06.

- [2] *Augur: the world's most accessible, no-limit betting platform*. <https://www.augur.net>. Accessed: 2020-02-16.
- [3] *Gnosis*. <https://gnosis.io/>. Accessed: 2020-04-16.
- [4] *Hivemind*. <https://hivmind.io/>. Accessed: 2020-04-16.
- [5] *Meet the 'Assassination Market' Creator Who's Crowdfunding Murder with Bitcoins*. <https://www.forbes.com/sites/andygreenberg/2013/11/18/meet-the-assassination-market-creator-whos-crowdfunding-murder-with-bitcoins/>. Accessed: 2020-06-06.
- [6] J. BERG AND T. RIETZ, *The iowa electronic markets: Stylized facts and open issues*, AEI- Brookings Joint Center for Regulatory Studies, (2006), pp. 142–169.
- [7] R. FREEMAN, S. LAHAIE, AND D. M. PENNOCK, *Crowdsourced Outcome Determination in Prediction Markets*, arXiv:1612.04885 [cs], (2016). arXiv: 1612.04885.
- [8] R. HANSON, *Logarithmic Market Scoring Rules for Modular Combinatorial Information Aggregation*, p. 12.
- [9] R. JURCA AND B. FALTINGS, *Incentives for expressing opinions in online polls*, in Proceedings of the 9th ACM conference on Electronic commerce - EC '08, Chicago, IL, USA, 2008, ACM Press, pp. 119–128.
- [10] —, *Incentives for Answering Hypothetical Questions*, in Workshop on Social Computing and User Generated Content - EC '11, 2011, p. 9.
- [11] C. KROER, M. DUDÍK, S. LAHAIE, AND S. BALAKRISHNAN, *Arbitrage-free combinatorial market making via integer programming*, in Proceedings of the 2016 ACM Conference on Economics and Computation, 2016, pp. 161–178.
- [12] Y. LIU, J. WANG, AND Y. CHEN, *Surrogate scoring rules*, in Proceedings of the 21st ACM Conference on Economics and Computation, EC '20, New York, NY, USA, 2020, Association for Computing Machinery, p. 853–871.
- [13] J. J. NAY, M. VAN DER LINDEN, AND J. M. GILLIGAN, *Betting and belief: prediction markets and attribution of climate change*, in 2016 Winter Simulation Conference (WSC), IEEE, 2016, pp. 1666–1677.
- [14] J. PETERSON, J. KRUG, M. ZOLTU, A. K. WILLIAMS, AND S. ALEXANDER, *Augur: a decentralized oracle and prediction market platform*, arXiv preprint arXiv:1501.01042, (2015).
- [15] B. SKIERA AND M. SPANN, *Internet-based virtual stock markets for business forecasting*, Management Science, 49 (2003), pp. 1310–1326.
- [16] J. WITKOWSKI, *Robust Peer Prediction Mechanisms*, PhD thesis, 2014.