

A Decentralised Prediction Market

CS907 Presentation

Thomas Archbold

April 21, 2020

Department of Computer Science, University of Warwick

Introduction

Prediction Markets allow you to trade on the outcomes of future events

Prediction Markets allow you to trade on the outcomes of future events

Market prices indicate how its traders feel about the outcome of an event

Prediction Markets allow you to trade on the outcomes of future events

Market prices indicate how its traders feel about the outcome of an event

- traders have different beliefs/knowledge informing their decisions

Prediction Markets allow you to trade on the outcomes of future events

Market prices indicate how its traders feel about the outcome of an event

- traders have different beliefs/knowledge informing their decisions
- provides means of aggregating information on topics of interest via crowdsourcing (“wisdom of the crowd”)

Prediction Markets allow you to trade on the outcomes of future events

Market prices indicate how its traders feel about the outcome of an event

- traders have different beliefs/knowledge informing their decisions
- provides means of aggregating information on topics of interest via crowdsourcing (“wisdom of the crowd”)

Securities are traded between 0 and 100%, and the payoff is contingent upon the outcome

Prediction Markets allow you to trade on the outcomes of future events

Market prices indicate how its traders feel about the outcome of an event

- traders have different beliefs/knowledge informing their decisions
- provides means of aggregating information on topics of interest via crowdsourcing (“wisdom of the crowd”)

Securities are traded between 0 and 100%, and the payoff is contingent upon the outcome

- typically, one share pays out \$1 in the event the outcome occurs, otherwise \$0

Traditional prediction markets are **centralised**

- trusted centre creates and closes market, manages transactions
- determines the outcome of the event and pays out accordingly

Traditional prediction markets are **centralised**

- trusted centre creates and closes market, manages transactions
- determines the outcome of the event and pays out accordingly

This restricts the types of bets that can be made

- must be explicitly offered by market maker/system
- high workload on trusted centre

A **decentralised market** allows anyone to create markets for events

- event outcome is determined by consensus among users (*arbiters*)

A **decentralised market** allows anyone to create markets for events

- event outcome is determined by consensus among users (*arbiters*)

Removes the need for a trusted centre BUT:

A **decentralised market** allows anyone to create markets for events

- event outcome is determined by consensus among users (*arbiters*)

Removes the need for a trusted centre BUT:

- bets may become ambiguous, since there is no central moderator

A **decentralised market** allows anyone to create markets for events

- event outcome is determined by consensus among users (*arbiters*)

Removes the need for a trusted centre BUT:

- bets may become ambiguous, since there is no central moderator
- arbiters may manipulate the outcome of the market

A **decentralised market** allows anyone to create markets for events

- event outcome is determined by consensus among users (*arbiters*)

Removes the need for a trusted centre BUT:

- bets may become ambiguous, since there is no central moderator
- arbiters may manipulate the outcome of the market
- users still need to bet based on their true beliefs

A **decentralised market** allows anyone to create markets for events

- event outcome is determined by consensus among users (*arbiters*)

Removes the need for a trusted centre BUT:

- bets may become ambiguous, since there is no central moderator
- arbiters may manipulate the outcome of the market
- users still need to bet based on their true beliefs

Freeman, Lahaie, and Pennock [1] introduced a mechanism that tackles these issues

Goal: implement a decentralised prediction market

Goal: implement a decentralised prediction market

In particular:

Goal: implement a decentralised prediction market

In particular:

- create a web-application where users may make and trade bets with virtual money

Goal: implement a decentralised prediction market

In particular:

- create a web-application where users may make and trade bets with virtual money
- place no restrictions on the bets that can be made

Goal: implement a decentralised prediction market

In particular:

- create a web-application where users may make and trade bets with virtual money
- place no restrictions on the bets that can be made
- outcome is determined by a group of arbiters (who may hold a stake in the market)

Goal: implement a decentralised prediction market

In particular:

- create a web-application where users may make and trade bets with virtual money
- place no restrictions on the bets that can be made
- outcome is determined by a group of arbiters (who may hold a stake in the market)

Why is this worthwhile?

- less restrictive way to crowdsource public opinion
- possibility for manipulation requires game-theoretic approach

Decentralised prediction markets based on cryptocurrencies:

Decentralised prediction markets based on cryptocurrencies:

- *Augur*
 - reputation system to protect against manipulation

Decentralised prediction markets based on cryptocurrencies:

- *Augur*
 - reputation system to protect against manipulation
- *Gnosis*
 - a distributed exchange for “tokens”
 - uses “trading rings” to trade the large numbers of illiquid bets

Decentralised prediction markets based on cryptocurrencies:

- *Augur*
 - reputation system to protect against manipulation
- *Gnosis*
 - a distributed exchange for “tokens”
 - uses “trading rings” to trade the large numbers of illiquid bets
- *Hivemind*

Decentralised prediction markets based on cryptocurrencies:

- *Augur*
 - reputation system to protect against manipulation
- *Gnosis*
 - a distributed exchange for “tokens”
 - uses “trading rings” to trade the large numbers of illiquid bets
- *Hivemind*

Predictalot was a centralised **combinatorial** prediction market

- combinatorial bids
- combinatorial outcomes

Design Overview

Market stage

- set up market for event X using a **market scoring rule** (MSR)
- agents trade in the market
- the market closes and trading stops

Arbitration stage

- each arbiter receives signal (e.g. reading the news) and reports the outcome
- arbiters are paired randomly and paid according to 1/prior mechanism
- outcome of the market is the fraction of arbiters reporting $X = 1$

Language: Lisp (SBCL)

Packages:

- Hunchentoot
- CL-WHO
- Parensript
- Mito

- Environment setup
- Skeleton code for mechanism
- Logarithmic Market Scoring Rule
- 1/prior payment mechanism

The web server is run on Hunchentoot

```
1      (setf *web-server* (make-instance 'hunchentoot:easy-acceptor 8080))  
2      (hunchentoot:start *web-server*)
```

HTML is generated using macros supplied by CL-WHO

```
1      (cl-who:with-html-output-to-string ...)
```

Macros are used define page templates

```
1      (defmacro define-url-fn ((name) &body body)  
2          ...)
```


We have two options to implement a prediction market:

We have two options to implement a prediction market:

- Continuous Double Auction (CDA)

We have two options to implement a prediction market:

- Continuous Double Auction (CDA)
- Market Scoring Rule (MSR)

We have two options to implement a prediction market:

- Continuous Double Auction (CDA)
- Market Scoring Rule (MSR)

In CDAs, an order book tracks bids and asks

- if a bid matches an ask, the trade is executed
- used in traditional stock markets

Problem: low liquidity

- with few participants, there is no guarantee any bid will match an ask

Problem: low liquidity

- with few participants, there is no guarantee any bid will match an ask

Solution: an *automated market maker* which always assumes the opposite side of a trade

- participants are always able to trade
- BUT the “house” is likely to lose money

Scoring rules measure the accuracy of probabilistic predictions

- we use the **Logarithmic Market Scoring Rule** [2]

$$C_b(q) = b \cdot \log(1 + e^{q/b}) \quad (1)$$

where:

- q is the quantity (number of shares)
- b is an arbitrary constant (“liquidity” parameter)
 - smaller value \Rightarrow higher price responsiveness

An agent wishing to buy $q' - q$ shares of a security pays $C_b(q') - C_b(q)$

- suppose $b = 10$ and there are 10 outstanding shares
- Alice wishes to buy 7 shares $\Rightarrow q' = 17$
- she pays $C_{10}(17) - C_{10}(10) = \5.54

An agent wishing to buy $q' - q$ shares of a security pays $C_b(q') - C_b(q)$

- suppose $b = 10$ and there are 10 outstanding shares
- Alice wishes to buy 7 shares $\Rightarrow q' = 17$
- she pays $C_{10}(17) - C_{10}(10) = \5.54

Share price is calculated with:

$$p(q) = \frac{e^{q/b}}{1 + e^{q/b}} \quad (2)$$

An agent wishing to buy $q' - q$ shares of a security pays $C_b(q') - C_b(q)$

- suppose $b = 10$ and there are 10 outstanding shares
- Alice wishes to buy 7 shares $\Rightarrow q' = 17$
- she pays $C_{10}(17) - C_{10}(10) = \5.54

Share price is calculated with:

$$p(q) = \frac{e^{q/b}}{1 + e^{q/b}} \quad (2)$$

$p(q)$ is not used to charge users, only to calculate closing price of the market

The outcome of an event is not binary – rather, it takes value $\hat{X} \in [0, 1]$

- corresponds to the proportion that reported that the event occurred

The outcome of an event is not binary – rather, it takes value $\hat{X} \in [0, 1]$

- corresponds to the proportion that reported that the event occurred

Two nice properties:

The outcome of an event is not binary – rather, it takes value $\hat{X} \in [0, 1]$

- corresponds to the proportion that reported that the event occurred

Two nice properties:

1. each question has well-defined, unambiguous outcome

The outcome of an event is not binary – rather, it takes value $\hat{X} \in [0, 1]$

- corresponds to the proportion that reported that the event occurred

Two nice properties:

1. each question has well-defined, unambiguous outcome
2. limits the influence a single arbiter can have on the value of a security

Trading fees imposed on the worst-case loss incurred by an agent

1. raises funds to pay the arbiters
2. bounds security prices away from 0 and 1

Trading fees imposed on the worst-case loss incurred by an agent

1. raises funds to pay the arbiters
2. bounds security prices away from 0 and 1

For *liquidation* transactions, no fees are imposed

- users are selling shares they own, or buying back those they have sold

Trading fees imposed on the worst-case loss incurred by an agent

1. raises funds to pay the arbiters
2. bounds security prices away from 0 and 1

For *liquidation* transactions, no fees are imposed

- users are selling shares they own, or buying back those they have sold

Otherwise, the user is charged

- fp for a security bought at price p
- $f(1 - p)$ for security sold at price p

Arbiters need to be incentivised to act truthfully, since they may hold a stake in the market

This is achieved using (a modified) 1/prior mechanism:

$$u(\hat{x}_i, \hat{x}_j) = \begin{cases} k\mu & \text{if } \hat{x}_i = \hat{x}_j = 1 \\ k(1 - \mu) & \text{if } \hat{x}_i = \hat{x}_j = 0 \\ 0 & \text{otherwise} \end{cases} \quad (3)$$

where:

- k is an arbitrary constant
- μ is the prior probability that an agent receives a positive signal

The mechanism uses the following parameters:

- b – liquidity
- k – to pay the arbiters
- f – trading fee
- δ – “update strength”, measures how strongly correlated signals are between arbiters

Once the system is running we will need to tune these to suit the market:

- number of participants, ambiguity of question, maximum budget of a user, etc.

Project Management

Project has changed slightly from that of the research proposal

- still looking to incorporate some aspects of it

Project has changed slightly from that of the research proposal

- still looking to incorporate some aspects of it

Much of the coming term will be spent revising for exams

- basic functionality not far off working, but needs to work in a web application
- bulk of the project to be done over the summer

Once the skeleton code is finished it will need integrating into a *website*

The next major tasks will be to:

- design and implement the database to store information on all users and markets
- creating the pages through which users will interact with the system
- tuning the parameters of the mechanism

Beyond this is the (stretch) goal of computing relationships between markets

- how does a bet placed on one security affect the odds of another?

Timetable – out with the old...




Figure 1: Timetable as of the research proposal submission

	02/17	02/24	03/02	03/09	03/16	03/23	03/30	04/06	04/13	04/20	04/27	05/04	05/11	05/18	05/25	06/01	06/08	06/15	06/22
Exam period																			
Specification																			
Presentation																			
Environment setup																			
Web server setup																			
Market Scoring Rule																			
Market stage																			
Arbitration stage																			
Database setup/design																			
Web interface																			
Parameter tuning																			
Prediction Markets literature																			

Figure 2: Revised timetable

Demonstration

Thank you for listening!

References

- [1] Rupert Freeman, Sebastien Lahaie, and David M. Pennock. Crowdsourced Outcome Determination in Prediction Markets. *arXiv:1612.04885 [cs]*, December 2016. URL <http://arxiv.org/abs/1612.04885>. arXiv: 1612.04885.
- [2] Robin Hanson. Logarithmic Market Scoring Rules for Modular Combinatorial Information Aggregation. page 12.